

Trabalho 3

Descrição do trabalho:

<http://webserver2.tecgraf.puc-rio.br/~mgattass/visao/trb/T3.html> (<http://webserver2.tecgraf.puc-rio.br/~mgattass/visao/trb/T3.html>).

Aluno: Daniel da Silva Costa

E-mail: danieldasilvacosta@gmail.com

Imports

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
import time
from keras.datasets import mnist
from keras.utils import np_utils
from sklearn import metrics
import pandas as pd

import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
```

Using the Numpy

Based on <https://www.youtube.com/watch?v=M0jIQzHo5Jg> (<https://www.youtube.com/watch?v=M0jIQzHo5Jg>).

Utility functions

In [2]:

```
def to_categorical(x, n_col=None):
    """ One hot encoding function """
    if not n_col:
        n_col = np.amax(x) + 1

    one_hot = np.zeros((x.shape[0], n_col))
    one_hot[np.arange(x.shape[0]), x] = 1
    return one_hot
```

In [3]:

```
def accuracy(y_true, y_pred):
    """ Returns accuracy """
    return np.sum(y_true == y_pred, axis = 0) / len(y_true)
```

In [4]:

```
def batch_loader(X, y = None, batch_size=64):  
    """ Generates batches for training """  
    n_samples = X.shape[0]  
    for i in np.arange(0, n_samples, batch_size):  
        begin, end = i, min(i + batch_size, n_samples)  
        if y is not None:  
            yield X[begin:end], y[begin: end]  
        else:  
            yield X[begin:end]
```

Loss

In [5]:

```
class CrossEntropy():  
    def __init__(self): pass  
  
    def loss(self, y, p):  
        p = np.clip(p, 1e-15, 1- 1e-15)  
        return -y*np.log(p) - (1 - y) * np.log(1- p)  
  
    def gradient(self, y, p):  
        p = np.clip(p, 1e-15, 1- 1e-15)  
        return -(y/p) + (1 - y) / (1 - p)
```

Activation Functions

In [6]:

```
class LeakyReLU():  
    def __init__(self, alpha = 0.2):  
        self.alpha = alpha  
  
    def __call__(self, x):  
        return self.activation(x)  
  
    def activation(self, x):  
        return np.where(x >= 0, x, self.alpha * x)  
  
    def gradient(self, x):  
        return np.where(x >= 0, 1, self.alpha)
```

In [7]:

```

class Softmax():
    def __init__(self): pass

    def __call__(self, x):
        return self.activation(x)

    def activation(self, x):
        e_x = np.exp(x - np.max(x, axis = -1, keepdims=True))
        return e_x / np.sum(e_x, axis=-1, keepdims = True)

    def gradient(self, x):
        # Error was in our softmax
        p = self.activation(x)
        return p * (1 - p)

```

Layers

class Linear

Fully Connected Layer

In [8]:

```

class Linear():
    def __init__(self, n_in, n_out, name="linear"):
        limit = 1 / np.sqrt(n_in)
        self.W = np.random.uniform(-limit, limit, (n_in, n_out))
        self.b = np.zeros((1, n_out)) # Biases
        self.input = None
        self.output = None
        self.name = name

    def forward(self, x):
        self.input = x
        self.output = np.dot(self.input, self.W) + self.b # Wx + b
        return self.output

    def backward(self, output_gradient, learning_rate = 0.01):
        input_gradient = np.dot(output_gradient, self.W.T)
        weights_gradients = np.dot(self.input.T, output_gradient) # Calculate the weights error
        bias_gradients = np.sum(output_gradient, axis=0).reshape((1, -1))

        # Usually, we would allow an optimizer function to update the weights
        # but here, we just use simple SGD
        self.W -= learning_rate * weights_gradients
        self.b -= learning_rate * bias_gradients
        # We updated the weights below
        return input_gradient

    def __call__(self, x):
        return self.forward(x)

```

class Activation

In [9]:

```

class Activation():
    def __init__(self, activation, name="activation"):
        self.activation = activation
        self.gradient = activation.gradient
        self.input = None
        self.output = None
        self.name = name

    def forward(self, x):
        self.input = x
        self.output = self.activation(x)
        return self.output

    def backward(self, output_error, learning_rate = 0.01):
        return self.gradient(self.input) * output_error

    def __call__(self, x):
        return self.forward(x)

```

Network

In [10]:

```

class Network():
    def __init__(self, input_dim, output_dim, learning_rate=0.01):
        # input_dim = 784, output_dim = 10 for mnist
        self.layers = [
            Linear(input_dim, 256, name="input"),
            Activation(LeakyReLU(), name="relu1"),
            Linear(256, 128, name="input"),
            Activation(LeakyReLU(), name="relu2"),
            Linear(128, output_dim, name="output"),
            Activation(Softmax(), name="softmax")
        ]
        self.learning_rate = learning_rate

    def forward(self, x):
        for layer in self.layers:
            x = layer(x)
        return x

    def backward(self, loss_grad):
        for layer in reversed(self.layers):
            loss_grad = layer.backward(loss_grad, self.learning_rate)
        # Iterating backwards through the layers

    def __call__(self, x):
        return self.forward(x)

```


Loading the data

In [11]:

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [12]:

```
fig, axes = plt.subplots( 1, 1, figsize=( 10, 5 ) )  
  
axes.set_title( f'Label: {y_train[0]}' )  
axes.imshow( X_train[0], cmap='gray' )  
  
plt.show()
```



In [13]:

```
# X_train[0]
```

Depicting examples

In [14]:

```

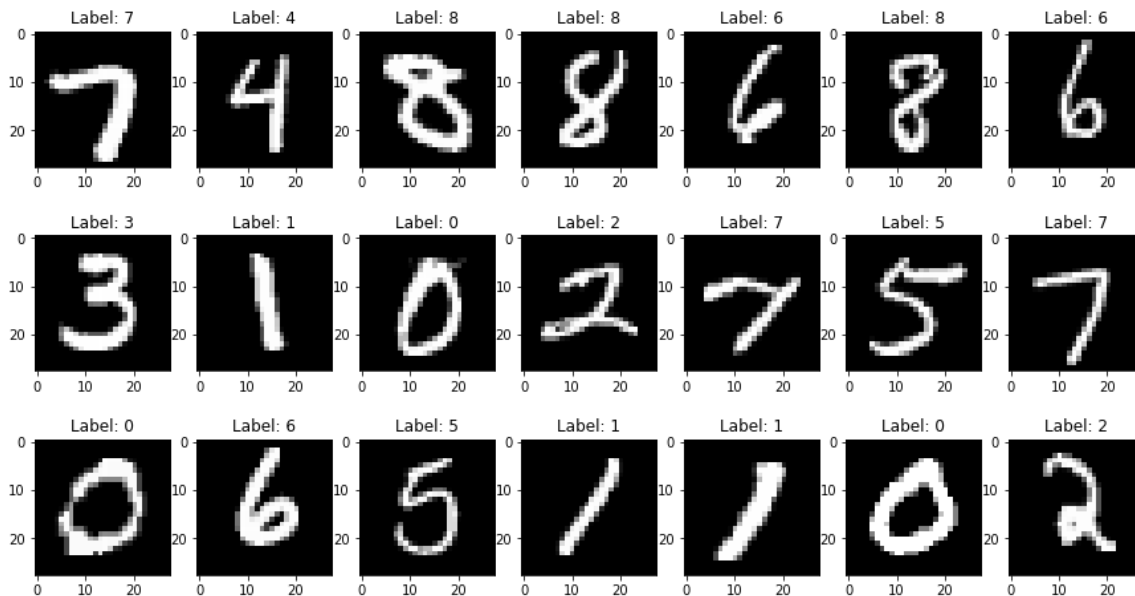
rows = 3
columns = 7

fig, axes = plt.subplots( rows, columns, figsize=( 15, 8 ) )

for i in range(rows):
    for j in range(columns):
        random_number = int(np.random.rand() * X_train.shape[0])
        axes[i][j].set_title( f'Label: {y_train[random_number]}' )
        axes[i][j].imshow( X_train[random_number], cmap='gray' )

plt.show()

```



Preprocessing Data

In [15]:

```

y_train, y_test = to_categorical(y_train.astype("int")), to_categorical(y_test.astype(
"int"))
X_train, X_test = X_train / 255.0, X_test / 255.0

```

In [16]:

```

X_train, X_test = X_train.reshape(-1, 28*28), X_test.reshape(-1, 28*28)
X_train.shape, X_test.shape

```

Out[16]:

```
((60000, 784), (10000, 784))
```

Training

In [17]:

```
n_input_dim = 28*28 # 784
n_out = 10 # 10 classes
```

In [18]:

```
criterion = CrossEntropy()
model = Network(n_input_dim, n_out, learning_rate=1e-3)
```

In [19]:

```
EPOCHS = 5
# batch_size = 1
batch_size = 128
# batch_size = X_train.shape[0]
```

In [20]:

```
%%time

for epoch in range(EPOCHS):
    loss = []
    acc = []
    for x_batch, y_batch in batch_loader(X_train, y_train, batch_size):
        # Forward pass
        out = model(x_batch)

        # Loss - for display
        loss.append(np.mean(criterion.loss(y_batch, out)))

        # Accuracy - For display
        acc.append(accuracy(np.argmax(y_batch, axis=1), np.argmax(out, axis=1)))

        # Calculate gradient of Loss
        error = criterion.gradient(y_batch, out)

        # Backpropagation
        model.backward(error)

    print(f"Epoch {epoch + 1}, Loss: {np.mean(loss):.5}, Acc: {np.mean(acc):.3}")
```

```
Epoch 1, Loss: 0.095875, Acc: 0.836
Epoch 2, Loss: 0.043966, Acc: 0.925
Epoch 3, Loss: 0.033638, Acc: 0.944
Epoch 4, Loss: 0.026969, Acc: 0.955
Epoch 5, Loss: 0.022378, Acc: 0.963
CPU times: user 28.7 s, sys: 11.1 s, total: 39.8 s
Wall time: 29.3 s
```

Testing

Accuracy

In [21]:

```
%time  
  
out = model(X_test)  
accuracy_value = accuracy(np.argmax(y_test, axis=1), np.argmax(out, axis=1))  
print( f'{{(accuracy_value * 100):.3}%}' )
```

CPU times: user 3 μ s, sys: 1 μ s, total: 4 μ s
Wall time: 7.39 μ s
96.0%

Error Rate

"o valor da taxa de erro (calculada como o número de todas as previsões incorretas dividido pelo número total do conjunto de dados. A melhor taxa de erro é 0, enquanto a pior é 1.)"

In [22]:

```
y_predicted = list(np.argmax(out, axis=1))  
y_actual = list(np.argmax(y_test, axis=1))  
  
for i, y_pred in enumerate(y_predicted):  
    print('pred:', y_predicted[i], '\ttrue:', y_actual[i])  
  
    if i > 20:  
        break
```

pred: 7	true: 7
pred: 2	true: 2
pred: 1	true: 1
pred: 0	true: 0
pred: 4	true: 4
pred: 1	true: 1
pred: 4	true: 4
pred: 9	true: 9
pred: 6	true: 5
pred: 9	true: 9
pred: 0	true: 0
pred: 6	true: 6
pred: 9	true: 9
pred: 0	true: 0
pred: 1	true: 1
pred: 5	true: 5
pred: 9	true: 9
pred: 7	true: 7
pred: 3	true: 3
pred: 4	true: 4
pred: 9	true: 9
pred: 6	true: 6

In [23]:

```

num_errors = 0
for i, y_pred in enumerate(y_predicted):
    if y_predicted[i] != y_actual[i]:
        # print('pred:', y_predicted[i], '\ttrue:', y_actual[i])
        num_errors += 1

print( f'num_errors: {num_errors}' )

```

num_errors: 402

In [24]:

```

error_rate = num_errors / X_train.shape[0]
print( f'error_rate: {error_rate:.3}' )

```

error_rate: 0.0067

In [25]:

```

error_rate_percentage = num_errors / X_train.shape[0] * 100
print( f'error_rate_percentage: {error_rate_percentage:.2}%' )

```

error_rate_percentage: 0.67%

Confusion Matrix

In [26]:

```

confusion_matrix = metrics.confusion_matrix(y_actual, y_predicted)

```

In [27]:

```

# based on https://www.statology.org/confusion-matrix-python/

y_actual = pd.Series(y_actual, name='Actual')
y_predicted = pd.Series(y_predicted, name='Predicted')

#create confusion matrix
print(pd.crosstab(y_actual, y_predicted))


```

Predicted \ Actual	0	1	2	3	4	5	6	7	8	9
0	968	0	0	2	0	5	3	0	1	1
1	0	1123	2	1	0	1	3	1	4	0
2	8	7	981	7	7	2	5	4	11	0
3	1	0	2	961	2	21	0	7	10	6
4	1	0	5	0	948	0	7	1	3	17
5	7	2	0	3	2	862	7	0	5	4
6	10	3	0	0	3	15	924	1	2	0
7	0	14	14	7	4	1	0	948	2	38
8	4	4	2	9	7	13	6	3	923	3
9	6	8	1	6	18	6	1	1	2	960

In [28]:

```
# based on https://www.w3schools.com/python/python_ml_confusion_matrix.asp

cm_display = metrics.ConfusionMatrixDisplay(
    confusion_matrix = confusion_matrix,
    display_labels = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
cm_display.plot()
plt.show()
```



Using PyTorch

Based on <https://www.youtube.com/watch?v=oPhxf2fXHkQ&list=PLqnsIRFeH2UrcDBWF5mfPGpqQDSta6VK4&index=13>

<https://www.youtube.com/watch?v=oPhxf2fXHkQ&list=PLqnsIRFeH2UrcDBWF5mfPGpqQDSta6VK4&index=13>

<https://www.youtube.com/watch?v=oPhxf2fXHkQ&list=PLqnsIRFeH2UrcDBWF5mfPGpqQDSta6VK4&index=13>

<https://www.youtube.com/watch?v=oPhxf2fXHkQ&list=PLqnsIRFeH2UrcDBWF5mfPGpqQDSta6VK4&index=13>

Device configuration

In [29]:

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
device
```

Out[29]:

```
device(type='cuda')
```

Hyper-parameters

In [30]:

```
input_size = 784 # 28x28
hidden_size = 500
num_classes = 10
num_epochs = 5
batch_size = 256
learning_rate = 0.001
```

MNIST dataset

In [31]:

```
train_dataset = torchvision.datasets.MNIST(root='./data',
                                           train=True,
                                           transform=transforms.ToTensor(),
                                           download=True)

test_dataset = torchvision.datasets.MNIST(root='./data',
                                           train=False,
                                           transform=transforms.ToTensor())
```

Data loader

Creating the batches of the training set and test set.

In [32]:

```
train_loader = torch.utils.data.DataLoader(dataset=train_dataset,
                                           batch_size=batch_size,
                                           shuffle=True)

test_loader = torch.utils.data.DataLoader(dataset=test_dataset,
                                           batch_size=batch_size,
                                           shuffle=False)
```

In [33]:

```
train_loader
```

Out[33]:

```
<torch.utils.data.dataloader.DataLoader at 0x7fc1ff3dda90>
```

In [34]:

```
# list(enumerate(train_loader))[0]
```

Depicting examples

From test set


In [35]:

```

examples = iter(test_loader)
example_data, example_targets = examples.next()

for i in range(6):
    plt.subplot(2,3,i+1)
    plt.imshow(example_data[i][0], cmap='gray')
plt.show()

```



Neural Network

In [36]:

```

# Fully connected neural network with one hidden layer
class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super(NeuralNet, self).__init__()
        self.input_size = input_size
        self.l1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.l2 = nn.Linear(hidden_size, num_classes)

    def forward(self, x):
        out = self.l1(x)
        out = self.relu(out)
        out = self.l2(out)
        # no activation and no softmax at the end
        return out

```

In [37]:

```

model = NeuralNet(input_size, hidden_size, num_classes).to(device)

```

Loss and optimizer

In [38]:

```
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
```

Training

In [39]:

```
%%time

n_total_steps = len(train_loader)
for epoch in range(num_epochs):
    for i, (images, labels) in enumerate(train_loader):
        # origin shape: [100, 1, 28, 28]
        # resized: [100, 784]
        images = images.reshape(-1, 28*28).to(device)
        labels = labels.to(device)

        # Forward pass
        outputs = model(images)
        loss = criterion(outputs, labels)

        # Backward and optimize
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        if (i+1) % 100 == 0:
            print (f'Epoch [{epoch+1}/{num_epochs}], Step [{i+1}/{n_total_steps}], Loss: {loss.item():.4f}')
```

```
Epoch [1/5], Step [100/235], Loss: 0.2875
Epoch [1/5], Step [200/235], Loss: 0.2020
Epoch [2/5], Step [100/235], Loss: 0.1686
Epoch [2/5], Step [200/235], Loss: 0.1448
Epoch [3/5], Step [100/235], Loss: 0.1988
Epoch [3/5], Step [200/235], Loss: 0.0796
Epoch [4/5], Step [100/235], Loss: 0.0650
Epoch [4/5], Step [200/235], Loss: 0.0768
Epoch [5/5], Step [100/235], Loss: 0.0524
Epoch [5/5], Step [200/235], Loss: 0.0713
CPU times: user 23.3 s, sys: 145 ms, total: 23.4 s
Wall time: 23.4 s
```

Testing

In [40]:

```
with torch.no_grad():
    n_correct = 0
    n_samples = 0
    for images, labels in test_loader:
        images = images.reshape(-1, 28*28).to(device)
        labels = labels.to(device)
        outputs = model(images)
        # max returns (value ,index)
        _, predicted = torch.max(outputs.data, 1)
        n_samples += labels.size(0)
        n_correct += (predicted == labels).sum().item()

    acc = 100.0 * n_correct / n_samples
    print(f'Accuracy of the network on the {n_samples} test images: {acc} %')
```

Accuracy of the network on the 10000 test images: 97.63 %

Exporting to HTML file

Changing Google Drive folder.

In [41]:

```
from google.colab import drive
drive.mount('/content/drive/', force_remount=True)
```

Mounted at /content/drive/

In [42]:

```
!cd "/content/drive/MyDrive/Doutorado/Disciplinas/[2022.2] [PUC-Rio] Visão Computacional - Professor: Marcelo Gattass/Trabalhos/Trabalho 3"
```

In [43]:

```
!pwd
```

/content

In [45]:

```
%%time
```

```
!jupyter nbconvert --to html ./T3_DanielCosta.ipynb
```